

Othniel Obasi

AI Strategist | ML/LLM Product Architect | NLP & Fraud Systems Expert

London, UK • +44-790-106-2935 • othnielobasi@gmail.com

- 6+ yrs shipping ML & LLM products across fintech, SaaS, utilities.
- Fraud & safety systems that cut loss **25 %+** and escalations **30 %**.
- Bridges strategy ↔ research ↔ engineering at C-suite and code levels.
- Proven track record in ML/LLM building and deployment, explainable AI pipelines, and safety-hardened LLM applications.

Key Skills & Tools

Python

FastAPI

PyTorch

scikit-learn

LSTM

Neo4j

OpenTelemetry

SHAP/LIME/DiCE

RAG

RLHF

Kafka

Pinecone

Docker

GCP

GitHub Actions

Data Analysis

Explainable AI

Fraud Detection

LLM Safety

Cross-functional Collaboration

Business Analysis

Product Ownership

Road Mapping

Experience

AI Strategist - Bezaleel Solutions

2023 – Present

Driving AI system design, LLM safety, and observability across mission-critical products.

- **NL→SQL Analyst Agent:** Designed and deployed an autonomous natural-language-to-SQL pipeline with 95 % parsing accuracy and 100 % execution success.
- **Integrated LLM-as-a-judge** (92 % clarity rating) and instrumented OpenTelemetry tracing into Arize Phoenix.
- **SaaS Support Bot Hardening:** Defined multi-tier safety rails for a billing support bot, blocked 85 %+ of 50+ red-team vectors, maintained 100 % citation compliance, and reduced human escalations by 30 %.

Technical Product Manager - ML - Creditville MFBank

Led ML-driven product development for fraud, risk, and customer intelligence in a digital finance ecosystem.

2022 – 2023

- Delivered 3 ML services: fraud detection, sentiment analytics, and auto loan decisioning — yielding double-digit efficiency gains.
- Built a modular explainability engine (LIME, DiCE) with visualizations and audit logs, integrated into a FastAPI fraud pipeline, cutting analyst triage time by 40 %.
- Deployed an NLP insights API using Word2Vec, LDA, and BERT summarization, reducing manual review by 50 %.
- Designed a risk-feature engine with real-time monitoring, reducing loan defaults by 20 %.

Lead, AI & Analytics - PHED (Power Holding Company)

Pioneered PHED's first AI/ML programs to modernize grid integrity, forecasting, and operational automation.

2020 – 2022

- Built real-time energy-fraud model, reducing undetected loss by 25 %.
- Delivered cloud-based forecasting with 95 % confidence; accelerated budgeting cycle by 2 weeks.
- Automated workflows (Python, Power Automate), cutting manual work by 30 %.

- Built executive dashboards (Excel, Tableau, Power BI) processing 2 M+ records/month.

Senior Business Analyst & Product Owner - *PHED* 2016 – 2020

Orchestrated process redesign and digital transformation pilots to prepare for enterprise AI adoption.

- Mapped 80+ processes via Lean redesign, cutting hand-offs and incidents by 30 %.
- Coordinated IoT pilots across 42 units; co-designed PoC dashboards with tech partners.
- Founded a Change Management Committee; trained 200+ staff on new tools.

Additional Roles (Part-Time)

Founding Technical Lead - AI/ML - Tech Genius Solutions

Driving prototype-stage AI infrastructure from fraud detection to SaaS-ready deployment across real-time systems. 2024 – Present

- Architected real-time fraud system (Kafka → PyTorch → Neo4j), boosting ring recall by 28 %.
- Built explainability dashboards with audit-ready outputs.
- Automated retraining loops; reduced missed fraud by 17 % and enabled future SaaS scaling.
- Integrated adaptive biometric/MFA verifications, cutting false positives by 35 %.

LLM Training & Evaluation Specialist - Outlier AI &

Invisible Technologies Jan 2024 – Jan 2025 | Remote
Executed high-throughput LLM evaluation and RLHF for two San Francisco, US AI firms.

- Evaluated LLM outputs for bias, factuality, and safety; flagged 150+ issues.
- Created 300+ prompt-completion pairs with RLHF feedback, improving win rate by 11 %.
- Audited model behavior and applied prompt tuning + data augmentation for safety.

Selected Projects - POC/Research

Fraud Detection via Logic Tensor Networks

Python • PyTorch • LTN • scikit-learn

Achieved 93.1 % MLP accuracy (5-fold CV); added fuzzy-logic constraints using LTN.

Final logic-consistency score: 0.94; optimized CPU performance for 500-epoch runs.

LexSynergy AI - Legal Decision System

Legal-BERT • LSTM • Neo4j • SHAP • RAG • FastAPI

90 % accuracy in case outcome prediction via hybrid neuro-symbolic model.

Embedded neural outputs into Neo4j graph with Cypher-based retrieval.

Applied SHAP/LIME for legal explainability on precedent- and statute-level insights.

Distilled & Quantised ML Service (MNIST Full Stack)

PyTorch • FastAPI • MLflow • W&B • PostgreSQL • Docker • GCP • GitHub Actions

Trained, distilled, and quantized a digit classifier achieving 98 % accuracy with sub-100 ms inference latency; tracked evaluation metrics using MLflow and W&B.

Integrated predictions with PostgreSQL for persistent storage and analytics.

Containerized with Docker, deployed on GCP, and automated CI/CD with GitHub Actions for zero-click updates.

Education & Certifications

- **M.Sc. Financial Engineering** – WorldQuant University, US
- **M.Sc. Applied AI & Data Science (Distinction)** – Southampton Solent University, UK
- Neo4j Certified Professional
- Applied Data Science – Scientific Computing & Python (Honors) – WorldQuant University

Last updated: Apr 2025